

Importing all the necessary libraries

```
In [72]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn import svm
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
```

Importing the dataset to understand the structure of the dataset

```
In [73]: diabetes_data = pd.read_csv('Diabetes_dataset.csv')
diabetes_data.head()
```

```
Out[73]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

```
In [74]: diabetes_data.tail()
```

```
Out[74]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
763	10	101	76	48	180	32.9	0.171	63	0
764	2	122	70	27	0	36.8	0.340	27	0
765	5	121	72	23	112	26.2	0.245	30	0
766	1	126	60	0	0	30.1	0.349	47	1
767	1	93	70	31	0	30.4	0.315	23	0

Number of rows and columns in the dataset

```
In [75]: diabetes_data.shape
print("The number of rows in the dataset are: ", diabetes_data.shape[0])
print("The number of columns in the dataset are: ", diabetes_data.shape[1])
```

The number of rows in the dataset are: 768
The number of columns in the dataset are: 9

Checking for any empty values in the dataset

```
In [76]: diabetes_data.isnull().sum()
```

```
Out[76]: Pregnancies      0
Glucose      0
BloodPressure 0
SkinThickness 0
Insulin      0
BMI          0
DiabetesPedigreeFunction 0
Age          0
Outcome      0
dtype: int64
```

Looking at the statistics of the dataset

```
In [77]: diabetes_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column              Non-Null Count  Dtype
---  -
0   Pregnancies         768 non-null    int64
1   Glucose             768 non-null    int64
2   BloodPressure       768 non-null    int64
3   SkinThickness       768 non-null    int64
4   Insulin             768 non-null    int64
5   BMI                 768 non-null    float64
6   DiabetesPedigreeFunction 768 non-null    float64
7   Age                 768 non-null    int64
8   Outcome             768 non-null    int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

```
In [78]: diabetes_data.describe()
```

```
Out[78]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578	0.471876	33.240885	0.348958
std	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.078000	21.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000	0.243750	24.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000	0.372500	29.000000	0.000000
75%	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	0.626250	41.000000	1.000000
max	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	2.420000	81.000000	1.000000

Total number of Predicted Non-Diabetic and Diabetic Patients

```
In [79]: diabetes_data['Outcome'].value_counts()
```

```
Out[79]: Outcome
0      580
1      268
Name: count, dtype: int64
```

```
In [80]: X = diabetes_data.drop(columns = 'Outcome', axis = 1)
Y = diabetes_data['Outcome']
```

```
In [81]: print(X)
```

```
Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin  BMI  \
0             6      148             72             35      0   33.6
1             1       85             66             29      0   26.6
2             8      183             64              0      0   23.3
3             1       89             66             23     94   28.1
4             0      137             40             35     168   43.1
..          ...     ...             ...           ...     ...     ...
763          10      101             76             48     180   32.9
764           2      122             70             27      0   36.8
765           5      121             72             23     112   26.2
766           1      126             60              0      0   30.1
767           1       93             70             31      0   30.4

DiabetesPedigreeFunction  Age
0             0.627      50
1             0.351      31
2             0.672      32
3             0.167      21
4             2.288      33
..          ...     ...
763          0.171      63
764          0.340      27
765          0.245      30
766          0.349      47
767          0.315      23
```

[768 rows x 8 columns]

```
In [82]: print(Y)
```

```
0      1
1      0
2      1
3      0
4      1
..     -
763     0
764     0
765     0
766     1
767     0
Name: Outcome, Length: 768, dtype: int64
```

Standardizing the input features to ensure data consistency, thereby enhancing the accuracy of our SVM model

```
In [83]: scaler = StandardScaler()
scaler.fit(X)
standardized_data = scaler.transform(X)
print(standardized_data)

[[ 0.63994726  0.84832379  0.14964075 ...  0.20401277  0.46849198
  1.4259954 ]
 [-0.84488505 -1.12339636 -0.16054575 ... -0.68442195 -0.36506078
 -0.19067191]
 [ 1.23388019  1.94372388 -0.26394125 ... -1.10325546  0.60439732
 -0.10558415]
 ...
 [ 0.34298088  0.00330087  0.14964075 ... -0.73518964 -0.68519336
 -0.27575966]
 [-0.84488505  0.1597866  -0.47073225 ... -0.24020459 -0.37110101
  1.17073215]
 [-0.84488505 -0.8730192  0.04624525 ... -0.20212881 -0.47378505
 -0.87137393]]
```

Building our svm model and feeding the standardized data

```
In [84]: X = standardized_data
Y = diabetes_data['Outcome']
print(X)
print(Y)

[[ 0.63994726  0.84832379  0.14964075 ...  0.20401277  0.46849198
  1.4259954 ]
 [-0.84488505 -1.12339636 -0.16054575 ... -0.68442195 -0.36506078
 -0.19067191]
 [ 1.23388019  1.94372388 -0.26394125 ... -1.10325546  0.60439732
 -0.10558415]
 ...
 [ 0.34298088  0.00330087  0.14964075 ... -0.73518964 -0.68519336
 -0.27575966]
 [-0.84488505  0.1597866  -0.47073225 ... -0.24020459 -0.37110101
  1.17073215]
 [-0.84488505 -0.8730192  0.04624525 ... -0.20212881 -0.47378505
 -0.87137393]]

0      1
1      0
2      1
3      0
4      1
..     -
763     0
764     0
765     0
766     1
767     0
Name: Outcome, Length: 768, dtype: int64
```

Splitting the data for training and testing

```
In [85]: X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = 0.2, stratify = Y, random_state = 4)
print(X.shape, X_train.shape, X_test.shape)
```

(768, 8) (614, 8) (154, 8)

Using Support Vector Classifier for classification of the data

```
In [86]: classifier = svm.SVC(kernel = 'linear')
classifier.fit(X_train, Y_train)
```

```
Out[86]: SVC(kernel='linear')
```

Testing the accuracy of our model on training data

```
In [87]: X_train_prediction = classifier.predict(X_train)
training_data_accuracy = accuracy_score(X_train_prediction, Y_train)
print("The accuracy of the training data is: ", training_data_accuracy)
print(confusion_matrix(X_train_prediction, Y_train))
```

The accuracy of the training data is: 0.7866449511400652
[[362 93]
 [38 121]]

Testing the accuracy of our model on testing data

```
In [88]: X_test_prediction = classifier.predict(X_test)
test_data_accuracy = accuracy_score(X_test_prediction, Y_test)
print("The accuracy of the test data is: ", test_data_accuracy)
print(confusion_matrix(X_test_prediction, Y_test))
```

The accuracy of the test data is: 0.7272727272727273
[[79 21]
 [21 33]]

Supplying one of the inputs to check the model

```
In [89]: input_data = (1,103,30,38,83,43.3,0.183,33)
input_data_as_numpy_array = np.asarray(input_data)
input_data_reshaped = input_data_as_numpy_array.reshape(1, -1)
std_data = scaler.transform(input_data_reshaped)
print(std_data)

[[-0.84488505 -0.56004775 -2.02166474  1.09545411  0.02778979  1.43512945
 -0.87244072 -0.0204964 ]]
```

C:\Users\91957\AppData\Roaming\Python\Python312\site-packages\sklearn\base.py:493: UserWarning: X does not have valid feature names, but StandardScaler was fitted with feature names
warnings.warn(

```
In [90]: prediction = classifier.predict(std_data)
print(prediction)
if prediction[0] == 0:
    print("The person is not diabetic")
else:
    print("The person is diabetic")

[0]
The person is not diabetic
```